

GAURAV RAKESH SINHA

sinha.g@northeastern.edu | +1-857-398-3073 | [LinkedIn](#) | [GitHub](#) | [Portfolio](#) | Santa Clara, CA

Education

Master of Science in Computer Software Engineering <i>Northeastern University</i> Boston, MA	Sep 2023 - Aug 2025
<u>Specialized Stream:</u> Web Design / User Experience, Cloud Computing, Theory & Applications of Gen-AI	
Bachelor of Engineering in Electronics and Telecommunication <i>Mumbai University</i> India	Sep 2017 - May 2021

Technical Skills

Programming Languages:	Java, Python, JavaScript, TypeScript, SQL
Web Technologies:	React, Next.js, Node.js, React Native, FastAPI, Jest, JUnit, Selenium, OAuth, gRPC, Expo, Figma
AI & Frameworks:	LangChain, Google ADK, FastMCP, LLMs, RAG, Fine-tuning, PyTorch, Pinecone, NLP
Databases & Monitoring:	PostgreSQL, MongoDB, Redis, Kafka, ELK stack, SonarQube, Grafana, Splunk
DevOps & Tools:	Git, GitHub, Docker, Kubernetes, AWS (EC2, ECS, S3, Lambda), Jenkins, JIRA, Ansible, Bash

Professional Experience

XNET Mobile <i>Irvine, CA</i>	<i>Full Stack Developer Intern</i>	Jan 2025 - May 2025
<ul style="list-style-type: none">Built REST APIs and customer management features using Next.js to accelerate slow onboarding process, adding OAuth 2.0 and JWT authentication system for secure login and integrating Solana Wallet functionality that streamlined account registration workflowContainerized services using Docker to ensure consistent deployment across environments, integrating Jenkins CI/CD pipeline that automated builds and deployments to AWS infrastructure and achieved 95% successful deployment rateDeployed containerized application to AWS ECS with auto-scaling to support reliable production infrastructure, configuring CloudWatch monitoring and alarms to alert the team about API errors and performance issues in real-time		
Tata Consultancy Services <i>Mumbai, India</i>	<i>Software Engineer</i>	Aug 2022 - Jul 2023
<ul style="list-style-type: none">Designed versioned REST APIs within Spring Boot microservices for credit card benefits platform, implementing v1/v2 strategy to support parallel legacy jQuery and React clients during migration, serving 10,000+ concurrent usersCreated search API with Elasticsearch integration after monitoring revealed 28% of users abandoning searches due to slow queries on 5,000+ product dataset, implementing full-text search with query optimization and result pagination that cut abandonment to 15%Established comprehensive JUnit test suites to address quality concerns and reduce bugs reaching production, achieving 90% code coverage across microservices components and strengthening software reliabilityConfigured PostgreSQL indexing and Redis caching layer to improve database query performance, reducing API response times by 55% on frequently accessed endpoints and decreasing database load for product catalog queries		
Tata Consultancy Services <i>Mumbai, India</i>	<i>Associate Software Engineer</i>	Jul 2021 - Jul 2022
<ul style="list-style-type: none">Engineered real-time Market Analytics dashboard with React and TypeScript to replace static charts frustrating traders, integrating WebSocket connections for live price updates that reduced data display latency from 350ms to 200msOptimized portfolio watchlist architecture using Redux Toolkit with normalized state structure to eliminate redundant API calls that degraded app performance, improving watchlist load time from 4.2 seconds to 2.8 secondsEnhanced trading platform accessibility to meet client's WCAG 2.1 compliance requirements, implementing screen reader support, keyboard navigation, and semantic HTML across core features that improved accessibility score from 23% to 85%		

Projects

<u>Presentation Generator</u>	(<i>Next.js, Redux, React DND, MongoDB, TypeScript</i>)
<ul style="list-style-type: none">Constructed a dynamic web application leveraging Next.js and React DND, enabling users to design and share custom animation-based presentations, increasing user engagement through interactive experiencesIntegrated Redux to store element states, enhancing application responsiveness and reducing state-related errors by 25%	
<u>MCP-Powered Zoo Tour Guide Agent</u>	(<i>FastMCP, Google ADK, LangChain, Docker, Google Cloud Run</i>)
<ul style="list-style-type: none">Built a tour guide system by deploying a secure FastMCP server and an ADK-based AI agent integrated with LangChain on Cloud Run, enabling the Gemini 2.5 Flash model to access external zoo data and Wikipedia knowledge via the MCP protocolContainerized Python microservices with Docker using FastMCP for tool APIs and Google ADK for reasoning workflows, implementing IAM authentication and Cloud Logging for secure, production-grade orchestration	
<u>Equity Lens</u> Founding Engineer	(<i>FastAPI, Express, LangChain, FastMCP, Pinecone, React/TypeScript</i>)
<ul style="list-style-type: none">Engineered an AI agent using LangChain and FastMCP that reads user portfolio data, retrieves historical earnings transcripts from Pinecone, and fetches live market data via Yahoo Finance to answer investment questions in a single contextual responseDeveloped a full-stack platform with React frontend and Express backend, connected to a FastAPI AI service powered by a fine-tuned GPT-4o-mini trained on financial Q&A data for accurate, domain-specific responses	